

Contact

7207894708 (Mobile)
pradeepkumarbsf7@gmail.com

www.linkedin.com/in/pradeepkumar-snowflake-developer/ (LinkedIn)

Top Skills

Snowflake
Snowpark
SQL

Pradeep Kumar

Senior Snowflake Developer | Snowflake | Snowpark | Snow-SQL | Snowpipe | SQL | Python | DBT | PySpark | Azure | AWS | Power BI | Tableau | DevOps | Jira | ServiceNow | Scala | R | ML | DL | NLP | AI
Pune, Maharashtra, India

Experience

ZS

Senior Snowflake Developer
December 2024 - Present (8 months)
Pune, Maharashtra, India

Designed a scalable, cloud-native Snowflake data warehouse architecture to support ZS's data-driven consulting solutions across healthcare, pharma & life sciences domains.

Built robust data models & schemas optimized for large-scale analytical workloads, focusing on sales, marketing, and patient-level data integration.

Developed end-to-end ETL/ELT pipelines using Snowflake, DBT & orchestration tools to ingest data from disparate sources such as CRM systems, third-party pharma data providers & flat files.

Integrated multiple data streams including physician data, product sales, prescription data & market research to build unified data marts for downstream analytics.

Implemented transformation logic using complex SQL, Snowpark (Python), and DBT to clean, normalize, & enrich raw datasets into structured, analytics-ready tables.

Optimized Snowflake workloads with techniques like clustering keys, multi-cluster warehouses, caching, and query pruning to ensure high performance & cost-effective execution.

Enforced healthcare compliance & data governance practices by managing RBAC (Role-Based Access Control), masking policies, and secure data sharing in Snowflake.

Worked closely with analytics and BI teams to enable self-service reporting and dashboarding using Power BI and other visualization tools.

Supported campaign analysis, sales rep effectiveness, targeting strategy, and patient engagement analytics through custom-built Snowflake data solutions.

Collaborated with cross-functional teams including data scientists, product managers & cloud architects to deliver scalable and reusable Snowflake components.

Maintained version-controlled Snowflake code & CI/CD pipelines using Azure DevOps and Git.

Enabled historical tracking & slowly changing dimensions (SCDs) for healthcare provider, prescriber, and product hierarchy datasets.

Performed data validation, profiling & troubleshooting of production jobs, ensuring data accuracy, completeness & timely delivery for business users.

CareOregon

Senior Snowflake Developer

June 2024 - August 2024 (3 months)

Pune, Maharashtra, India

Design scalable and optimized Snowflake data warehouse architecture to store, process, and manage healthcare data.

Develop efficient data models for healthcare-specific data, ensuring compliance with data warehousing best practices.

Implement data warehouse solutions that support clinical & business needs.

Set up Snowflake data pipelines for loading data from multiple sources (like healthcare systems, flat files, etc.) using ETL, ELT tools.

Integrate data from electronic health records (EHR), claims data, provider data & other healthcare datasets.

Optimize data ingestion processes for performance & cost efficiency.

Develop & maintain ETL/ELT workflows in Snowflake to transform raw healthcare data into meaningful datasets.

Write complex SQL queries to transform & aggregate healthcare data. Ensure data quality, consistency & integrity throughout the transformation processes.

Optimize query performance & warehouse usage by tuning the Snowflake environment, including clustering, partitioning & efficient use of virtual warehouses.

Monitor & troubleshoot slow-running queries & other performance issues.

Ensure that all data handling complies with healthcare data privacy regulations.

Implement data security protocols within Snowflake, including role-based access control & data encryption.

Work with Security tools or other tools to secure sensitive healthcare data.

Collaborate with business analysts, healthcare providers & other stakeholders to understand data requirements & translate them into Snowflake solutions.

Assist in developing healthcare dashboards & reports for clinical & financial performance analysis.

Analyze healthcare data patterns to support business decisions related to patient care, resource allocation & financial forecasting.

Perform data exploration & validation to ensure data is accurate & up-to-date for healthcare reporting.

Implement data retention policies & manage slowly changing dimensions (SCDs) in healthcare provider data.

Broadridge

Senior Snowflake Developer
February 2024 - April 2024 (3 months)
Pune, Maharashtra, India

Designing, developing, and managing the data warehouse architecture within Snowflake to support Broadridge's financial data.

Collaborating with data architects and business analysts to ensure that the data warehouse meets business requirements and reporting needs.

Implementing ETL/ELT processes for data ingestion, transformation, and storage from various data sources, such as relational databases and flat files.

Optimizing Snowflake SQL queries for performance and efficiency, especially for large financial datasets.

Implementing partitioning, clustering, and indexing strategies to improve query performance.

Managing data partitioning, clustering, and table optimization for better query execution times.

Implementing data quality checks and ensuring that the data is accurate, clean, and consistent.

Managing schema design for new tables and supporting the star or snowflake schema for better reporting and analysis.

Developing data pipelines to load data from various financial systems into Snowflake using ETL, ELT tools.

Setting up processes for incremental data loading and real-time data integration as required by the financial systems.

Working with business stakeholders, data engineers, and reporting teams to ensure smooth data availability for analytics and reporting.

Providing technical expertise to junior developers and helping them resolve issues or optimize their Snowflake-related work.

Participating in code reviews, design sessions, and collaborating on best practices for Snowflake development.

Ensuring that the Snowflake environment is secure and complies with financial industry standards, such as data encryption, access control, and data masking.

Managing role-based access controls (RBAC) to ensure data security across different financial departments and ensuring compliance with regulatory standards.

Creating and maintaining technical documentation for ETL/ELT workflows and processes within Snowflake.

UBS

Snowflake Developer

April 2023 - July 2023 (4 months)

Navi Mumbai, Maharashtra, India

Design and implement robust data pipelines using Snowflake to support the UBS banking data warehouse.

Develop scalable data models (star schema, snowflake schema) that align with UBS's business objectives.

Ingest data from various banking source systems (databases, flat files, APIs) into Snowflake using ETL, ELT tools or custom SQL scripts.

Optimize the Snowflake environment, focusing on partitioning, clustering, and materialized views to improve query performance.

Design and build ETL/ELT processes to load data into Snowflake from multiple UBS banking sources.

Use SQL and Snowflake-specific features (e.g., streams, tasks, time travel) to automate data transformations.

Analyze and optimize complex SQL queries, ensuring fast query performance for large-scale banking data.

Implement strategies to reduce Snowflake's compute costs by optimizing resource utilization.

Configure role-based access controls (RBAC), encryption, and data masking features to secure sensitive data (e.g., financial transactions, customer information).

Set up auditing and monitoring processes to track data usage and adherence to compliance.

Develop strategies for ensuring data quality, including validation processes, error handling, and anomaly detection in ETL pipelines.

Perform unit testing and collaborate with QA teams for end-to-end testing of Snowflake solutions.

Work closely with business analysts and data scientists to understand UBS's data needs, translating business requirements into technical solutions.

Provide technical leadership to ensure data architecture aligns with banking objectives (e.g., risk analysis, financial reporting).

Establish and enforce best practices for Snowflake development, including version control (Git), CI/CD workflows, and agile development methodologies. Collaborate with cloud architects and DevOps teams to integrate Snowflake with other UBS systems.

Mercedes-Benz India

Senior Data Engineer

February 2022 - February 2023 (1 year 1 month)

Bengaluru, Karnataka, India

Design and implement scalable data warehouse solutions using Snowflake to support manufacturing operations and analytics.

Develop optimized Snowflake schemas (e.g., star and snowflake schemas) for manufacturing data to ensure efficient query performance.

Collaborate with data architects to design data models tailored to specific business use cases in the manufacturing domain.

Develop ETL/ELT pipelines to ingest data from various manufacturing systems (such as MES, ERP, and IoT sensors) into Snowflake.

Integrate Snowflake with data ingestion ETL, ELT tools like or other platforms for seamless data extraction and loading.

Work with real-time and batch data processing from flat files, relational databases, and APIs.

Optimize Snowflake queries and storage for fast data retrieval, particularly for large-scale manufacturing data sets.

Implement partitioning, clustering, and pruning strategies in Snowflake to enhance performance and minimize costs.

Monitor and resolve performance issues in ETL/ELT pipelines, ensuring data quality and efficiency.

Use Snowflake's capabilities (e.g., SnowSQL, Snowpipe, UDFs) to perform data transformations for reporting and analytics.

Collaborate with data scientists and analysts to transform raw manufacturing data into useful metrics and KPIs, such as production efficiency, machine downtime, and product quality.

Develop stored procedures and user-defined functions in Snowflake to automate complex data transformations.

Maintain version control of SQL scripts and data models to ensure reproducibility and collaboration across teams.

Act as a liaison between business users and technical teams, ensuring data warehouse solutions align with manufacturing goals.

Troubleshoot data discrepancies, data flow errors, and other issues in real-time, ensuring minimal disruption to manufacturing operations.

Research and recommend new tools or processes for better data handling and analytics in the manufacturing context.

Accenture

Data Engineer

June 2018 - February 2022 (3 years 9 months)

Mumbai, Maharashtra, India

Design and build an optimized Snowflake data warehouse for retail data, including sales, inventory, customer, and product data.

Collaborate with business analysts and data architects to define data models that cater to the specific needs of the retail industry, such as seasonality, demand forecasting, and product categorization.

Develop retail-specific data marts for various business units like marketing, supply chain, and finance.

Create and manage ETL/ELT pipelines using Snowflake and related ETL, ELT tools.

Extract data from multiple sources such as Point of Sale (POS) systems, ERP platforms (like SAP), and eCommerce databases.

Implement real-time data loading with Snowpipe for use cases like live inventory tracking, customer purchase trends, and retail analytics.

Write advanced SQL queries for transforming raw retail data into insights related to sales trends, customer behavior, and product performance.

Develop transformation logic to calculate key retail metrics such as sales per store, inventory turnover, product returns, gross margins, and customer lifetime value.

Implement stored procedures for repetitive and complex data transformation logic, such as SKU-level data aggregations or seasonal adjustment calculations.

Integrate data from retail-specific platforms (e.g., CRM systems, ERP systems, supply chain systems) into Snowflake for centralized reporting.

Ensure data consistency and accuracy across retail domains, including inventory management, supply chain logistics, product catalog, and pricing systems.

Optimize Snowflake virtual warehouses for large-scale retail data processing, especially for high-volume data like transactional data, inventory updates, and customer purchase histories.

Enable efficient data access for retail reporting tools such as Power BI, Tableau, or Looker, to deliver insights like store performance comparisons, product sales breakdowns, and inventory levels.

Education

Jawaharlal Nehru Technological University Kakinada (JNTUK)
Bachelor of Technology - BTech, Electrical, Electronics and Communications
Engineering · (August 2015 - May 2018)